

ASTR8150/PHYS8150

Optimization

Fabien Baron

Georgia State University

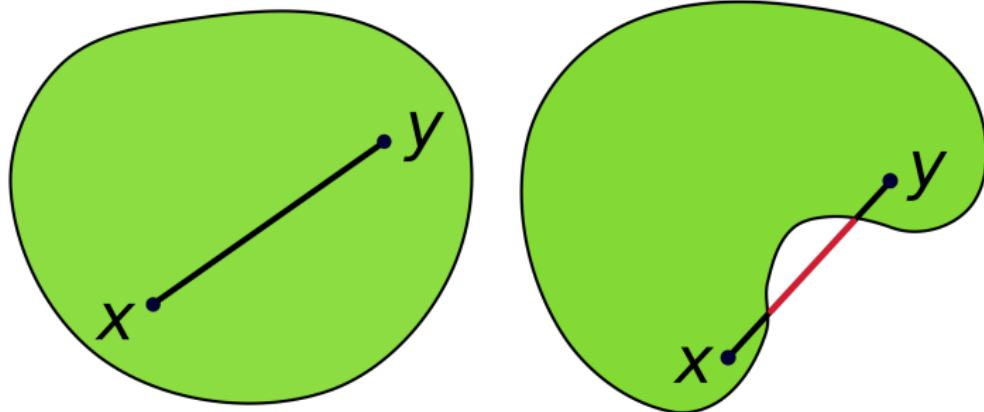
fbaron@gsu.edu

Fall 2025

Optimization Difficulty and Function Structure

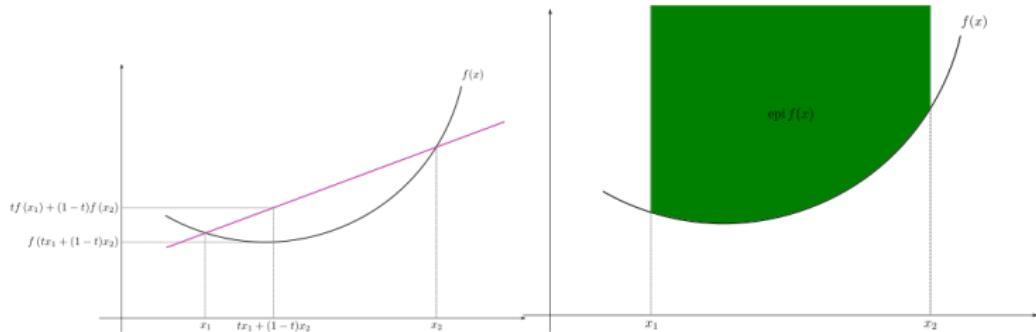
- Inference in a Bayesian framework will most often require the optimization of an objective function (e.g. minimization of the negative log-likelihood)
- The optimization difficulty depends on two properties:
 - ① **Convexity of the function:** will straight move within parameter space lead to complex behavior, such as local minima? Does every local minimum equal the global minimum?
 - ② **Smoothness of the function:** is there a gradient? Is it continuous and bounded?

Convexity of a set



- In a convex set, for every pair of points within the region, every point on the straight line segment that joins the pair of points is also within the region.
- A set which is hollow or has an indent, for example, a crescent shape, is not convex.

Convexity of a function



- A real-valued function is called convex if the set of points on or above the graph of the function (epigraph) is a convex set.
- For a twice differentiable function of a single variable, if the second derivative is always greater than or equal to zero for its entire domain then the function is convex. Examples: $f(x) = x^2$ or $f(x) = e^x$
- Jensen's inequality: if X is a convex set and $f : X \rightarrow \mathbb{R}$, f is convex if:

$$\forall x_1, x_2 \in X, \forall t \in [0, 1] : f(tx_1 + (1 - t)x_2) \leq tf(x_1) + (1 - t)f(x_2).$$

Smoothness of a function

- The smoothness of a function is a property measured by the number of derivatives it has which are continuous. A smooth function is a function that has derivatives of all orders everywhere in its domain.
- The function $f(x) = |x|^k$ is continuous and k times differentiable at all x . But at $x = 0$ it is not $(k + 1)$ times differentiable.
- The norms ℓ_2 , ℓ_1 and pseudo-norm ℓ_0 are used in regularization. ℓ_2 is convex, differentiable and smooth. ℓ_1 is convex, differentiable but nonsmooth. ℓ_0 is non-convex and nonsmooth.

Optimization of Convex and Smooth Functions

- Typical form: quadratic or log-sum-exp functions.
- Gradient is Lipschitz continuous:

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$$

- Guarantees:
 - Any local minimum is global.
 - Gradient descent converges linearly with suitable step size.
- Example: least-squares regression.

Optimization of Non-Convex but Smooth Functions

- Many local minima and saddle points.
- Gradient is well-behaved (differentiable, continuous), but:
 - Optimization can get trapped in local minima.
 - Second-order methods (Hessian) can help identify saddles.
- Example: neural network loss surfaces.
- Techniques:
 - Random initialization, momentum, stochasticity.
 - Trust-region or adaptive methods (Adam, LBFGS).

Optimization of Convex but Non-Smooth Functions

- Example: ℓ_1 -regularized objectives, such as classic total variation.
- Gradient may not exist everywhere.
- Use of *subgradients* or proximal operators (later in this chapter).
- Guarantees:
 - Still globally convex.
 - Convergence is typically sublinear ($<$ linear $<$ quadratic)
- Methods:
 - Subgradient descent, bundle methods, proximal gradient (ADMM).

Optimization of Non-Convex and Non-Smooth Functions

- Hardest class of problems ("NP-hard")
- Examples:
 - Compressed sensing with ℓ_0 , sparse deep learning, robust estimators, combinatorial losses.
- No general guarantees of convergence or global optimality.
- Strategies:
 - Initialization heuristics and regularization.
 - Relaxation ($\ell_0 \rightarrow \ell_1$) or smoothing of objectives.
 - Stochastic (e.g., simulated annealing) or global (some NLOpt methods) searches.

Derivative-free optimization methods

- **Derivative-free** optimization methods do not require information about the gradient to work
- These include most MCMC optimization methods (except for Hamiltonian Monte Carlo).
- Non-MCMC derivative free methods are ill-suited to optimize more than 3 parameters at once due to their slowness.
- Among the most popular local optimizer is **Nelder–Mead** method (aka downhill simplex method or amoeba method), which moves points of a polytope of $n + 1$ vertices in n -parameter dimensions via reflection, contraction, and expansion steps.
- The NLOpt library provides mostly derivative-free algorithms, some of them for global optimization.

Calculating gradients - Overview

- The gradient of the objective function with respect to the parameters to optimize gives the local slope of the function.
- Its usage will speed up optimization & machine learning codes significantly when compared to derivative-free methods.
- But first we need to calculate it
- Three main approaches (also used for Jacobian calculation):
 - ① Analytic differentiation
 - ② Numerical differentiation
 - ③ Automatic differentiation (AD)

Analytic Differentiation

- Compute derivatives symbolically from known formulas.
- Produces exact expressions (e.g. $\frac{d}{dx} \sin x = \cos x$).
- Pros:
 - Exact derivatives.
- Cons:
 - Tedious or intractable for complex programs.
 - Requires algebraic manipulation and simplification.

Numerical Differentiation

- Approximates derivative using finite differences:

$$f'(x) \approx \frac{f(x + h) - f(x)}{h}$$

- Pros:
 - Simple to implement.
- Cons:
 - Sensitive to step size h (round-off and truncation errors).
 - Computationally expensive for many parameters (requires multiple function calls).

Automatic Differentiation (AD)

- Computes exact derivatives using the chain rule at the level of elementary operations.
- Express program as a sequence of primitive operations:

$$x_1 = x, \quad x_2 = \sin(x_1), \quad x_3 = x_1 \cdot x_2, \text{ etc.}$$

- AD propagates derivatives through this computational graph.
- Two main modes:
 - Forward mode (tangent propagation)
 - Backward mode (adjoint or reverse accumulation)

Forward Mode AD

- For each intermediate variable v_i , compute both its value and its derivative \dot{v}_i .
- Using the chain rule:

$$\dot{v}_i = \sum_j \frac{\partial v_i}{\partial v_j} \dot{v}_j$$

- Efficient when the number of inputs is small (many outputs).
- Example:

$$f(x) = \sin(x^2) \implies \begin{cases} v_1 = x^2, & \dot{v}_1 = 2x \\ v_2 = \sin(v_1), & \dot{v}_2 = \cos(v_1)\dot{v}_1 \end{cases}$$

Backward Mode AD

- Computes derivatives from outputs back to inputs.
- Store the computation graph during the forward evaluation.
- In reverse pass, accumulate sensitivities:

$$\bar{v}_i = \sum_j \frac{\partial v_j}{\partial v_i} \bar{v}_j$$

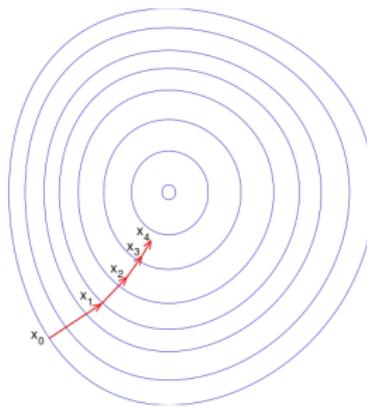
- Efficient when the number of outputs is small (many inputs), e.g. neural networks.
- Example: backpropagation for neural network.

Comparison of gradient obtention methods

Method	Exactness	Performance Cost	Scaling
Analytic	Exact	Lowest	Good but tedious
Numerical	Approximate	High (many evals)	Poor for many vars
AD (Forward)	Exact	Moderate	$\propto \# \text{inputs}$
AD (Backward)	Exact	Moderate	$\propto \# \text{outputs}$

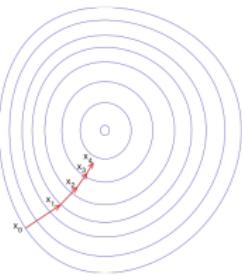
- Modern computer languages or libraries will provide forward and/or backward automatic differentiation with easy setup:
 - Python: Autograd, JAX, Pytorch/Tensorflow
 - Julia: ForwardDiff, ReverseDiff, Zygote, Enzyme, Mooncake
 - Matlab: ADiMat, MAD
 - Fortran and C/C++: Enzyme
 - IDL: nothing

Gradient descent (1)



- Gradient descent is based on the observation that if the multi-variable function $f(\mathbf{x})$ is defined differentiable in a neighborhood of a point \mathbf{x}_0 , then $f(\mathbf{x})$ decreases "fastest" if one goes from \mathbf{x}_0 in the direction of the negative gradient of f at \mathbf{x}_0 , $-\nabla f(\mathbf{x}_0)$.

Gradient descent (2)



- It follows that, if

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \alpha \nabla f(\mathbf{x}_n) \quad (1)$$

for α small enough, then $f(\mathbf{x}_n) \geq f(\mathbf{x}_{n+1})$. In other words, the term $\alpha \nabla f(\mathbf{x})$ is subtracted from \mathbf{x} because we want to move against the gradient, namely down toward the minimum.

- How can we choose α ?
- The Rosenbrock function $f(x_1, x_2) = (1 - x_1)^2 + 100(x_2 - x_1^2)^2$. has a narrow curved valley which contains the minimum. The bottom of the valley is very flat. Because of the curved flat valley the optimization is zig-zagging slowly with small stepsizes towards the minimum.

Steepest descent using line search

- **Inexact line search** consists in finding $\alpha_k \simeq \operatorname{argmin}_{\alpha \in \mathbb{R}_+} f(\mathbf{x}_k + \alpha \mathbf{d}_k)$
- The **line search** method is one of two basic iterative approaches to find a local minimum \mathbf{x}^* of an objective function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ using gradients. The other approach is **trust region**.

Algorithm 1 Steepest descent with line search

```
1: procedure STEEPEST DESCENT( $f, \mathbf{x}$ )
2:    $k = 0, \mathbf{x}_0$             $\triangleright$  Iteration counter + initial parameter guess
3:   while  $\|\nabla f(\mathbf{x}_k)\| > \epsilon$  do            $\triangleright \epsilon = \text{tolerance}$ 
4:      $\mathbf{d}_k = -\nabla f(\mathbf{x}_k)$             $\triangleright$  Descent direction = Steepest descent
5:      $\alpha_k \simeq \operatorname{argmin}_{\alpha \in \mathbb{R}_+} f(\mathbf{x}_k + \alpha \mathbf{d}_k)$             $\triangleright$  Line search
6:      $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$ 
7:      $k = k + 1$ 
8:   end while
9: end procedure
```

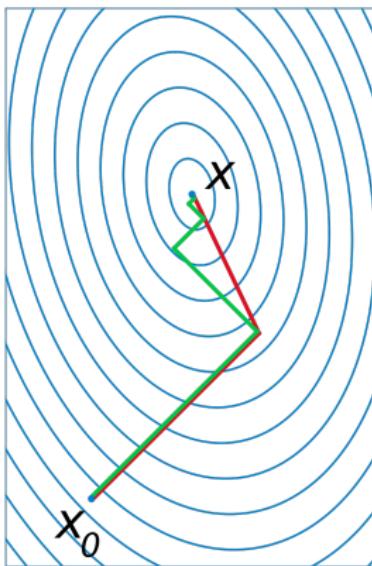
Nonlinear conjugate gradient methods

- Let's pose $\mathbf{g}_k = \nabla f(x_k)$
- The steepest descent direction was $\mathbf{d}_k = -\mathbf{g}_k$
- Conjugate directions differs from the steepest descent by attempting moves based on the history of the previous moves
- The descent direction for nonlinear conjugate gradient methods is

$$\mathbf{d}_{k+1} = -\mathbf{g}_{k+1} + \beta_k \mathbf{d}_k, \quad \mathbf{d}_0 = -\mathbf{g}_0 \quad (2)$$

- The variation of the gradient is measured by $\mathbf{y}_k = \mathbf{g}_{k+1} - \mathbf{g}_k$
- The Conjugate Gradient update parameter β_k can be updated with different formulas

Conjugate gradient: convergence



- A comparison of the linear convergence of simple gradient descent with optimal step size (in green) and the superlinear convergence of conjugate gradient (in red) for minimizing a quadratic function.

Nonlinear conjugate gradient methods

$$\beta_k^{HS} = \frac{\mathbf{g}_{k+1}^\top \mathbf{y}_k}{\mathbf{d}_k^\top \mathbf{y}_k} \quad (1952) \quad \text{in the original (linear) CG paper}$$

of Hestenes and Stiefel [59]

$$\beta_k^{FR} = \frac{\|\mathbf{g}_{k+1}\|^2}{\|\mathbf{g}_k\|^2} \quad (1964) \quad \text{first nonlinear CG method, proposed}$$

by Fletcher and Reeves [45]

$$\beta_k^D = \frac{\mathbf{g}_{k+1}^\top \nabla^2 f(\mathbf{x}_k) \mathbf{d}_k}{\mathbf{d}_k^\top \nabla^2 f(\mathbf{x}_k) \mathbf{d}_k} \quad (1967) \quad \text{proposed by Daniel [39], requires}$$

evaluation of the Hessian $\nabla^2 f(\mathbf{x})$

$$\beta_k^{PRP} = \frac{\mathbf{g}_{k+1}^\top \mathbf{y}_k}{\|\mathbf{g}_k\|^2} \quad (1969) \quad \text{proposed by Polak and Ribi  re [84]}$$

and by Polyak [85]

$$\beta_k^{CD} = \frac{\|\mathbf{g}_{k+1}\|^2}{-\mathbf{d}_k^\top \mathbf{g}_k} \quad (1987) \quad \text{proposed by Fletcher [44], CD}$$

stands for “Conjugate Descent”

$$\beta_k^{LS} = \frac{\mathbf{g}_{k+1}^\top \mathbf{y}_k}{-\mathbf{d}_k^\top \mathbf{g}_k} \quad (1991) \quad \text{proposed by Liu and Storey [67]}$$

$$\beta_k^{DY} = \frac{\|\mathbf{g}_{k+1}\|^2}{\mathbf{d}_k^\top \mathbf{y}_k} \quad (1999) \quad \text{proposed by Dai and Yuan [27]}$$

$$\beta_k^N = \left(\mathbf{y}_k - 2\mathbf{d}_k \frac{\|\mathbf{y}_k\|^2}{\mathbf{d}_k^\top \mathbf{y}_k} \right)^\top \frac{\mathbf{g}_{k+1}}{\mathbf{d}_k^\top \mathbf{y}_k} \quad (2005) \quad \text{proposed by Hager and Zhang [53]}$$

Newton optimization method

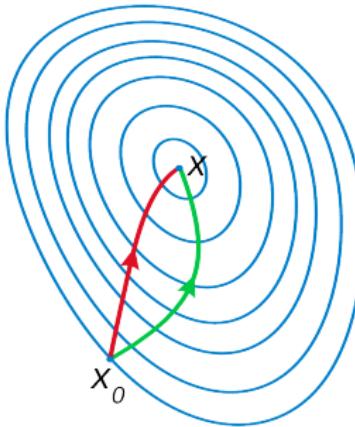


Figure: A comparison of gradient descent (green) and Newton's method (red) for minimizing a function (with small step sizes). Newton's method uses curvature information to take a more direct route.

- Hessian is used to exploit the curvature information

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \alpha [\mathbf{H}f(\mathbf{x}_n)]^{-1} \nabla f(\mathbf{x}_n) \quad (3)$$

- $\alpha \in (0, 1)$, with $\alpha = 1$ the exact form.

Newton-Raphson root-finding and eccentric anomaly

- Newton optimization method and Newton-Raphson's root finding methods are based on similar principles
- Newton-Raphson: $x_{n+1} = x_n - f(x_n)/f'(x_n)$
- Example: the mean anomaly is proportional to time it is an easily measured quantity for an orbiting body. Given the mean anomaly M , find the eccentric anomaly E and the orbital eccentricity e with Kepler's Equation:

$$M = E - e \sin E \tag{4}$$

Better than Newton: quasi-Newton methods

- Also known as variable metric methods, they avoid computing the Hessian then its inverse.
- The Broyden-Fletcher-Goldfarb-Shanno (BFGS) or Davidon-Fletcher-Powell (DFP) algorithms build iteratively approximations of $[\mathbf{H}f(\mathbf{x}_n)]^{-1}$.
- The most successful and well-known quasi-Newton method is the **Limited-memory BFGS (L-BFGS)** that approximates $[\mathbf{H}f(\mathbf{x}_n)]^{-1}\nabla f(\mathbf{x}_n)$ directly and thus can work on large scale problems (millions of variables).
- New gradient descent variants attempt to deal with non-smooth functions (subgradient and bundle method).
- There are variants that deal with constrained minimization (i.e. bounds on variables or linearly tied variables) such as L-BFGS-B. Further refinements led to the VMLM algorithm in OptimPack.

Trust-region method and Levenberg-Marquardt

- Consider the quadratic approximation of function f around x_0 :

$$q(\epsilon) \simeq f(x_0) + \nabla f(x_0)\epsilon + \frac{1}{2}\epsilon^T \nabla^2 f(x_0)\epsilon \quad (5)$$

- $q(\epsilon)$ has a close-form minimum.
- $q(\epsilon)$ remains a good approximation within a given radius, $\|\epsilon\|_2 < r^2$ defines the **trust region** radius r .
- The quadratic approximation predicts a certain reduction in the cost function, Δf_{pred} , which is compared to the true reduction $\Delta f_{\text{actual}} = f(x) - f(x + \epsilon)$. By looking at the ratio $\Delta f_{\text{pred}}/\Delta f_{\text{actual}}$ we can estimate the trust-region size at each iteration, jump to the closed-form minimum within the trust region, and iterate.
- The **Levenberg-Marquardt** algorithm (first published in 1944 by Kenneth Levenberg, rediscovered in 1963 by Donald Marquardt) uses the trust-region approach with conjugate-gradients and Gauss-Newton (Newton optimized for non-linear χ^2). Like conjugate gradient and Newton, these are local optimization codes.

Constrained minimization: the Lagrangian method

- **Constrained** minimization is minimization under equality or inequality constraints. **Bounded** optimization is a special case of constrained optimization where bounds are imposed on parameters (e.g. positivity, or variable within a range). In Bayesian terms, we're imposing a prior.
- A classic example is:

$$(\tilde{x}, \tilde{y}) = \underset{(x,y) \in \mathbb{R}^2}{\operatorname{argmin}} (x + y) \quad \text{s. t. } x^2 + y^2 = 1$$

- We pose $g(x, y) = x^2 + y^2 - 1$ and the Lagrangian is:

$$\begin{aligned}\mathcal{L}(x, y, \lambda) &= f(x, y) + \lambda \cdot g(x, y) \\ &= x + y + \lambda(x^2 + y^2 - 1).\end{aligned}$$

where λ is a Lagrange multiplier.

Constrained minimization: the Lagrangian method

- The gradient with respect to variables x, y and λ

$$\begin{aligned}\nabla_{x,y,\lambda} \mathcal{L}(x, y, \lambda) &= \left(\frac{\partial \mathcal{L}}{\partial x}, \frac{\partial \mathcal{L}}{\partial y}, \frac{\partial \mathcal{L}}{\partial \lambda} \right) \\ &= (1 + 2\lambda x, 1 + 2\lambda y, x^2 + y^2 - 1)\end{aligned}$$

and therefore:

$$\nabla_{x,y,\lambda} \mathcal{L}(x, y, \lambda) = 0 \quad \Leftrightarrow \quad \begin{cases} 1 + 2\lambda x = 0 \\ 1 + 2\lambda y = 0 \\ x^2 + y^2 - 1 = 0 \end{cases}$$

- Solution $x = y = -\frac{1}{2\lambda}$, $\lambda \neq 0$. Substituting into the last equation we get $\lambda = \pm \frac{1}{\sqrt{2}}$ which implies that the stationary points of \mathcal{L} are $\left(\frac{\sqrt{2}}{2}, \frac{\sqrt{2}}{2}, -\frac{1}{\sqrt{2}}\right)$, $\left(-\frac{\sqrt{2}}{2}, -\frac{\sqrt{2}}{2}, \frac{1}{\sqrt{2}}\right)$. And since $f\left(\frac{\sqrt{2}}{2}, \frac{\sqrt{2}}{2}\right) = \sqrt{2}$ and $f\left(-\frac{\sqrt{2}}{2}, -\frac{\sqrt{2}}{2}\right) = -\sqrt{2}$. the solution is found.

Constrained minimization: Half-quadratic splitting (1)

- Let's say we want to minimize:

$$\tilde{\mathbf{x}} = \operatorname{argmin}_{\mathbf{x}} \frac{1}{2} \|\mathbf{Hx} - \mathbf{y}\|_2^2 + \mu \Phi(\mathbf{x})$$

- If $\Phi(\mathbf{x}) = \|\mathbf{Wx} - \mathbf{w}\|_2^2$, Tikhonov gives us the solution
$$\mathbf{x} = (\mathbf{H}^\top \mathbf{H} + \mathbf{W}^\top \mathbf{W})^{-1}(\mathbf{H}^\top \mathbf{y} + \mathbf{W}^\top \mathbf{w})$$
- Other cases do not have closed-form solution, typical example is $\Phi(\mathbf{x}) = \ell_1(\mathbf{x}) = \|\mathbf{x}\|_1$.
- Splitting methods are methods that split the unconstrained problem into a constrained problem, using two different variables to represent the same one in different functions:

$$\begin{aligned}\tilde{\mathbf{x}} &= \operatorname{argmin}_{\mathbf{x}} \frac{1}{2} \|\mathbf{Hx} - \mathbf{y}\|_2^2 + \mu \Phi(\mathbf{z}) \quad s.t. \quad \mathbf{z} = \mathbf{x} \\ &= \operatorname{argmin}_{\mathbf{x}, \mathbf{z}} \frac{1}{2} \|\mathbf{Hx} - \mathbf{y}\|_2^2 + \mu \Phi(\mathbf{z}) + \frac{\rho}{2} \|\mathbf{z} - \mathbf{x}\|_2^2\end{aligned}$$

Constrained minimization: the two subproblems

- So we now want to minimize:

$$\mathcal{L}(\mathbf{x}, \mathbf{z}) = \frac{1}{2} \|\mathbf{H}\mathbf{x} - \mathbf{y}\|_2^2 + \mu\Phi(\mathbf{z}) + \frac{\rho}{2} \|\mathbf{z} - \mathbf{x}\|_2^2$$

- $\frac{\rho}{2} \|\mathbf{z} - \mathbf{x}\|_2^2$ is called an augmented term, and ρ is the augmented penalty hyperparameter.
- The **half-quadratic splitting method** solves iteratively (iteration variable = k) the problem with respect to \mathbf{x} , then \mathbf{z} :

$$\tilde{\mathbf{x}}^{k+1} = \operatorname{argmin}_{\mathbf{x}} \frac{1}{2} \|\mathbf{H}\mathbf{x} - \mathbf{y}\|_2^2 + \frac{\rho}{2} \|\tilde{\mathbf{z}}^k - \mathbf{x}\|_2^2 \quad \mathbf{x} \text{ sub-problem}$$

$$\tilde{\mathbf{z}}^{k+1} = \operatorname{argmin}_{\mathbf{z}} \frac{\rho}{2} \|\mathbf{z} - \tilde{\mathbf{x}}^{k+1}\|_2^2 + \mu\Phi(\mathbf{z}) \quad \mathbf{z} \text{ sub-problem}$$

and then increases ρ from initially low values to higher and higher ones.

Constrained minimization: analytical solutions exist

- Why is this easier than the original problem ?

$$\tilde{\mathbf{x}}^{k+1} = \operatorname{argmin}_{\mathbf{x}} \frac{1}{2} \|\mathbf{Hx} - \mathbf{y}\|_2^2 + \frac{\rho}{2} \left\| \tilde{\mathbf{z}}^k - \mathbf{x} \right\|_2^2 \quad \mathbf{x} \text{ sub-problem}$$

$$\tilde{\mathbf{z}}^{k+1} = \operatorname{argmin}_{\mathbf{z}} \frac{\rho}{2} \left\| \mathbf{z} - \tilde{\mathbf{x}}^{k+1} \right\|_2^2 + \mu \Phi(\mathbf{z}) \quad \mathbf{z} \text{ sub-problem}$$

- The \mathbf{x} sub-problem can be solved by classic Tikhonov.
- The \mathbf{z} sub-problem can be solved analytically for some functions, for which we know the solution of the problem:

$$\operatorname{prox}_f(\mathbf{x}) = \operatorname{argmin}_{\mathbf{z}} \left(\frac{1}{2} \|\mathbf{z} - \mathbf{x}\|_2^2 + f(\mathbf{z}) \right)$$

- This solution $\operatorname{prox}_f(\mathbf{x})$ is the **proximal operator** for the function f . At each point \mathbf{x} it finds a close-by local minimum of f .

Proximal operator for the ℓ_1 norm and positivity

- One can demonstrate that the proximal operator for ℓ_1 norm is:

$$\text{prox}_{\alpha \ell_1}(\mathbf{x}) = \underset{\mathbf{z}}{\operatorname{argmin}} \left(\frac{1}{2} \|\mathbf{z} - \mathbf{x}\|_2^2 + \alpha \ell_1(\mathbf{z}) \right) = \text{sign}(\mathbf{x}) \cdot \max(|\mathbf{x}| - \alpha, 0)$$

where \cdot is the Hadamard product.

- The proximal operator for positivity is the projection onto the positive set:

$$\text{prox}_{I_{\mathbb{R}^+}}(\mathbf{x}) = \underset{\mathbf{z} \in \mathbb{R}^{+n}}{\operatorname{argmin}} \left(\frac{1}{2} \|\mathbf{x} - \mathbf{z}\|_2^2 \right) = \max(\mathbf{x}, 0)$$

Half-quadratic splitting: beyond the 1:1 change of variable

- Half-quadratic splitting only involves analytical steps: the \mathbf{x} sub-problem is solved via Tikhonov and the \mathbf{z} sub-problem via proximal operators (provided it is known). Generalizing beyond $\mathbf{z} = \mathbf{x}$, we can also have more complex linear constraints under the form $\mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{z} + \mathbf{c} = 0$.
- How should we solve the total variation problem, i.e. minimize $\frac{1}{2} \|\mathbf{H}\mathbf{x} - \mathbf{y}\|_2^2 + \mu\ell_1(\nabla\mathbf{x})$?
- We only know the proximal operator for $\ell_1(\mathbf{z})$ and not for $\ell_1(\nabla\mathbf{z})$. So while we could pose $\mathbf{z} = \mathbf{x}$, we wouldn't know how to solve the \mathbf{z} sub-problem in closed form. However we can pose $\mathbf{z} = \nabla\mathbf{x}$, leading to:

$$\mathcal{L}(\mathbf{x}, \mathbf{z}) = \frac{1}{2} \|\mathbf{H}\mathbf{x} - \mathbf{y}\|_2^2 + \mu\ell_1(\mathbf{z}) + \frac{\rho}{2} \|\mathbf{z} - \nabla\mathbf{x}\|_2^2$$

Total variation solved via Half-quadratic splitting

$$\mathcal{L}(\mathbf{x}, \mathbf{z}) = \frac{1}{2} \|\mathbf{Hx} - \mathbf{y}\|_2^2 + \mu \ell_1(\mathbf{z}) + \frac{\rho}{2} \|\mathbf{z} - \nabla \mathbf{x}\|_2^2$$

- The \mathbf{x} sub-problem has the closed form Tikhonov solution:

$$\begin{aligned}\tilde{\mathbf{x}}^{k+1} &= \underset{\mathbf{x}}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{Hx} - \mathbf{y}\|_2^2 + \frac{\rho}{2} \|\tilde{\mathbf{z}}^k - \nabla \mathbf{x}\|_2^2 && \mathbf{x} \text{ sub-problem} \\ \implies \mathbf{H}^\top (\mathbf{Hx} - \mathbf{y}) - \rho \nabla^\top (\tilde{\mathbf{z}}^k - \nabla \mathbf{x}) &= 0 \\ \implies \tilde{\mathbf{x}}^{k+1} &= (\mathbf{H}^\top \mathbf{H} + \rho \nabla^\top \nabla)^{-1} (\mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k)\end{aligned}$$

- The \mathbf{z} sub-problem has a closed form proximal solution:

$$\begin{aligned}\tilde{\mathbf{z}}^{k+1} &= \underset{\mathbf{z}}{\operatorname{argmin}} \mu \ell_1(\mathbf{z}) + \frac{\rho}{2} \|\mathbf{z} - \nabla \tilde{\mathbf{x}}^{k+1}\|_2^2 && \mathbf{z} \text{ sub-problem} \\ &= \operatorname{prox}_{\frac{\mu}{\rho} \ell_1}(\nabla \tilde{\mathbf{x}}^{k+1}) = \operatorname{sign}(\nabla \tilde{\mathbf{x}}^{k+1}) \cdot \max(|\nabla \tilde{\mathbf{x}}^{k+1}| - \frac{\mu}{\rho}, 0)\end{aligned}$$

Making the x step faster: some tricks

- The x sub-problem is the slowest one since it involves:

$$\tilde{\mathbf{x}}^{k+1} = (\mathbf{H}^\top \mathbf{H} + \rho \nabla^\top \nabla)^{-1} (\mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k)$$

but matrix inversion is $\sim \mathcal{O}(N^3)$ process for a $N \times N$ matrix, and thus costly both in memory space and computing time. To work with larger images, tricks to speed up the inversion are used in practice.

- One is to use **sparse arrays** that only store the non-zero elements of \mathbf{H} and ∇ . The inversion and subsequent multiplication are then faster.
- Using an **orthogonal wavelet basis**, \mathbf{W} , as a sparsity basis instead of the spatial gradient ∇ since in this case $\mathbf{W}^\top \mathbf{W} = \alpha \mathbf{I}$.
- Another is to employ the backslash operator in Julia (or Matlab). We have $\mathbf{X}^{-1} \mathbf{Y} = \mathbf{X} \backslash \mathbf{Y}$, but the latter operation doesn't store the inverted matrix \mathbf{X}^{-1} ; instead it just temporarily stores the parts useful for its multiplication by \mathbf{Y} .

$$\tilde{\mathbf{x}}^{k+1} = (\mathbf{H}^\top \mathbf{H} + \rho \nabla^\top \nabla) \backslash (\mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k)$$

Making the x step faster: circulant matrices (1)

- It turns out $\nabla^\top \nabla$ and $\mathbf{H}^\top \mathbf{H}$ can be simplified if ∇ and \mathbf{H} are circulant matrices, or concatenation of circulant matrices. This is the case when ∇ is the spatial gradient, and when \mathbf{H} is a convolution (= if modeling the imaging done by an optical system).
- A circulant matrix \mathbf{C} takes the form:

$$\mathbf{C} = \begin{bmatrix} c_1 & c_n & \dots & c_3 & c_2 \\ c_2 & c_1 & c_n & & c_3 \\ \vdots & c_2 & c_1 & \ddots & \vdots \\ c_{n-1} & \vdots & \ddots & \ddots & c_n \\ c_n & c_{n-1} & \dots & c_2 & c_1 \end{bmatrix}.$$

- All circulant matrixes can be written as $\mathbf{C} = \mathbf{F}\mathbf{D}\mathbf{F}^\top$ where \mathbf{F} is the Fourier transform ($n^2 \times n^2$ if implemented via matrix operations) and $\mathbf{D} = \text{diag}(\mathbf{F}([c_1 \dots c_n]))$ is diagonal.

Making the x step faster: circulant matrices (2)

- We remind that $\mathbf{F}^\top \mathbf{F} = \mathbf{I}$ since $\mathbf{F}^\top = \mathbf{F}^{-1}$ and $(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top$
- Since $\mathbf{C} = \mathbf{FDF}^\top$, $\mathbf{C}^\top = ((\mathbf{FD})\mathbf{F}^\top)^\top = \mathbf{F}(\mathbf{FD})^\top = \mathbf{FD}^\top \mathbf{F}^\top$
- Consequently $\mathbf{C}^\top \mathbf{C} = \mathbf{FD}\cancel{\mathbf{F}^\top}\mathbf{FD}^\top \mathbf{F}^\top = \mathbf{FD}^2 \mathbf{F}^\top$ since \mathbf{D} is diagonal.
- The application relies on the fact that applying the Fourier transforms \mathbf{F} or \mathbf{F}^{-1} can be done with very fast FFT algorithms:

$$(\mathbf{H}^\top \mathbf{H} + \rho \nabla^\top \nabla) \tilde{\mathbf{x}}^{k+1} = \mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k$$

$$(\mathbf{F}\mathbf{D}_H^2 \mathbf{F}^\top + \rho \mathbf{F}\mathbf{D}_\nabla^2 \mathbf{F}^\top) \tilde{\mathbf{x}}^{k+1} = \mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k$$

$$\mathbf{F}(\mathbf{D}_H^2 + \rho \mathbf{D}_\nabla^2) \mathbf{F}^\top \tilde{\mathbf{x}}^{k+1} = \mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k$$

$$(\mathbf{D}_H^2 + \rho \mathbf{D}_\nabla^2) \mathbf{F}^\top \tilde{\mathbf{x}}^{k+1} = \mathbf{F}^{-1}(\mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k)$$

$$\mathbf{F}^\top \tilde{\mathbf{x}}^{k+1} = \frac{\mathbf{F}^{-1}(\mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k)}{\mathbf{D}_H^2 + \rho \mathbf{D}_\nabla^2}$$

$$\tilde{\mathbf{x}}^{k+1} = \mathbf{F} \left(\frac{\mathbf{F}^{-1}(\mathbf{H}^\top \mathbf{y} + \rho \nabla^\top \tilde{\mathbf{z}}^k)}{\mathbf{D}_H^2 + \rho \mathbf{D}_\nabla^2} \right)$$

Augmented Lagrangian methods

- Let's say we want to minimize $f(\mathbf{x})$ under a sets of constraints on \mathbf{x} . We saw we could express each of these constraints as $g_i(\mathbf{x}) = 0$. Then we just need to minimize the Lagrangian:

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \sum_i \lambda_i \cdot g_i(\mathbf{x})$$

where $\boldsymbol{\lambda}$ is a vector this time.

- The idea is to use an **augmented Lagrangian**:

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \sum_i \lambda_i \cdot g_i(\mathbf{x}) + \frac{\rho}{2} \sum_i |g_i(\mathbf{x})|^2$$

where the augmentation term $\sum_i |g_i(\mathbf{x})|^2$ is multiplied by penalty terms ρ .

- The **method of multipliers** solves this by iterating:

$$\mathbf{x}^{k+1} = \underset{\mathbf{x}}{\operatorname{argmin}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}_i^k)$$

$$\lambda_i^{k+1} = \lambda_i^k + \rho^k g_i(\mathbf{x}^k) \quad \forall i$$

Example: Method of multipliers for linear constraints

- Let's say we want to minimize $f(\mathbf{x})$ under a sets of linear constraints on \mathbf{x} , we can express them as $\mathbf{A}\mathbf{x} = \mathbf{b}$.
- The augmented Lagrangian is:

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \boldsymbol{\lambda}^\top (\mathbf{A}\mathbf{x} - \mathbf{b}) + \frac{\rho}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2^2$$

- The method of multipliers solves this by iterating:

$$\mathbf{x}^{k+1} = \operatorname{argmin}_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}^k)$$

$$\boldsymbol{\lambda}^{k+1} = \boldsymbol{\lambda}^k + \rho(\mathbf{A}\mathbf{x}^{k+1} - \mathbf{b})$$

Alternating Direction Method of Multipliers (ADMM)

- Putting all our previous knowledge together, using variable splitting, Lagrangian multipliers, and augmentation terms, we come with ADMM, a method to solve:

$$\operatorname{argmin}_{\mathbf{x}, \mathbf{z}} f(\mathbf{x}) + g(\mathbf{z}) \quad \text{s.t. } \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{z} + \mathbf{c} = 0$$

- ADMM solves this by posing:

$$\mathcal{L}(\mathbf{x}, \mathbf{z}, \boldsymbol{\lambda}) = f(\mathbf{x}) + g(\mathbf{z}) + \boldsymbol{\lambda}^\top (\mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{z} + \mathbf{c}) + \frac{\rho}{2} \|\mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{z} + \mathbf{c}\|_2^2$$

- And then we iterate:

$$\mathbf{x}^{k+1} = \operatorname{argmin}_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \mathbf{z}^k, \boldsymbol{\lambda}^k) = \operatorname{argmin}_{\mathbf{x}} f(\mathbf{x}) + \frac{\rho}{2} \left\| \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{z}^k + \mathbf{c} + \frac{\boldsymbol{\lambda}^k}{\rho} \right\|_2^2$$

$$\mathbf{z}^{k+1} = \operatorname{argmin}_{\mathbf{z}} \mathcal{L}(\mathbf{x}^{k+1}, \mathbf{z}, \boldsymbol{\lambda}^k) = \operatorname{argmin}_{\mathbf{z}} g(\mathbf{z}) + \frac{\rho}{2} \left\| \mathbf{A}\mathbf{x}^{k+1} + \mathbf{B}\mathbf{z} + \mathbf{c} + \frac{\boldsymbol{\lambda}^k}{\rho} \right\|_2^2$$

$$\boldsymbol{\lambda}^{k+1} = \boldsymbol{\lambda}^k + \rho(\mathbf{A}\mathbf{x}^{k+1} + \mathbf{B}\mathbf{z}^{k+1} + \mathbf{c})$$

Scaled ADMM

- For convenience we often find in paper the scaled form of ADMM, using the lagragian multiplier $\eta = \lambda/\rho$

$$\mathbf{x}^{k+1} = \operatorname{argmin}_{\mathbf{x}} f(\mathbf{x}) + \frac{\rho}{2} \left\| \mathbf{Ax} + \mathbf{Bz}^k + \mathbf{c} + \boldsymbol{\eta}^k \right\|_2^2$$

$$\mathbf{z}^{k+1} = \operatorname{argmin}_{\mathbf{z}} g(\mathbf{z}) + \frac{\rho}{2} \left\| \mathbf{Ax}^{k+1} + \mathbf{Bz} + \mathbf{c} + \boldsymbol{\eta}^k \right\|_2^2$$

$$\boldsymbol{\eta}^{k+1} = \boldsymbol{\eta}^k + (\mathbf{Ax}^{k+1} + \mathbf{Bz}^{k+1} + \mathbf{c})$$

- η is just the running sum of residuals.

ADMM residuals and convergence

- Primal residual: $\mathbf{r}^{k+1} = \mathbf{A}\mathbf{x}^{k+1} + \mathbf{B}\mathbf{z}^{k+1} + \mathbf{c}$, i.e. checks if the relation between x and z works.
- Dual residual: $\mathbf{s} = \rho \mathbf{A}^\top \mathbf{B}(\mathbf{z}^{k+1} - \mathbf{z}^k)$, i.e. if z has converged
- Convergence properties: global convergence of the iterative procedure to a local optimum is guaranteed. Slow (linear) rate of convergence, slower than gradient-based methods in theory. But the steps can be made very fast.
- Adaptive ADMM modifies ρ based on heuristics such as residual norm balancing or spectral methods.

ADMM: application to TV regularization with Gaussian likelihood

- Like with half quadratic, we pose $f(\mathbf{x}) = \frac{1}{2} \|\mathbf{Hx} - \mathbf{y}\|_2^2$, $g(\mathbf{z}) = \ell_1(\mathbf{z})$ and $\mathbf{z} = \nabla \mathbf{x}$, i.e. $\mathbf{A} = \nabla$, $\mathbf{B} = -\mathbf{I}$ and $\mathbf{c} = \mathbf{0}$.

$$\mathbf{x}^{k+1} = \operatorname{argmin}_{\mathbf{x}} f(\mathbf{x}) + \frac{\rho}{2} \left\| \nabla \mathbf{x} - \mathbf{z}^k + \boldsymbol{\eta}^k \right\|_2^2 \rightarrow \text{Tikhonov}$$

$$\mathbf{z}^{k+1} = \operatorname{argmin}_{\mathbf{z}} g(\mathbf{z}) + \frac{\rho}{2} \left\| \nabla \mathbf{x}^{k+1} - \mathbf{z} + \boldsymbol{\eta}^k \right\|_2^2 \rightarrow \text{Proximal operator for } g$$

$$\boldsymbol{\eta}^{k+1} = \boldsymbol{\eta}^k + \nabla \mathbf{x}^{k+1} - \mathbf{z}^{k+1}$$

These steps are very similar to Half Quadratic, with just the addition of η in the squared norms.

Consensus ADMM

- Consensus ADMM is solving $\operatorname{argmin}_{\mathbf{x}} \sum_{i=1}^N f_i(\mathbf{x})$ by recasting it as:

$$\operatorname{argmin}_{\mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{x}} \sum_{i=1}^N f_i(\mathbf{x}_i) \quad \text{s.t. } \mathbf{x}_i = \mathbf{x}, \quad \forall i.$$

- Used to add more regularization terms. But more importantly if the log-likelihood can be split (into different wavelengths, epochs, etc.), such as is often the case with χ^2 then ADMM allows for **efficient parallelization** of large-scale problems.
- Reminder: k = iteration index, i = term index.

$$\mathbf{x}_i^{k+1} = \operatorname{argmin}_{\mathbf{x}_i} f_i(\mathbf{x}_i) + \frac{\rho}{2} \left\| \mathbf{x}_i - \mathbf{x}^k + \boldsymbol{\eta}^k \right\|_2^2 \rightarrow \text{computed in parallel}$$

$$\mathbf{x}^{k+1} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i^{k+1} + \boldsymbol{\eta}^k) \rightarrow \text{consensus step, fast}$$

$$\boldsymbol{\eta}^{k+1} = \boldsymbol{\eta}^k + \mathbf{x}_i^{k+1} - \mathbf{x}^{k+1} \rightarrow \text{fast}$$

Interesting ADMM papers

- Deblurring with Poisson noise:
<https://ieeexplore.ieee.org/document/5492199>
- Deblurring with unknown boundaries:
<https://ieeexplore.ieee.org/document/6738120>