ASTR8150/PHYS8150
Signal Analysis
Fourier Transforms, Gaussian Processes

Fabien Baron

Georgia State University

baron@chara.gsu.edu

Fall 2019

DFT: Discrete Fourier Transform

- The Discrete Fourier Transform (DFT) transforms a sequence of N complex numbers $\{\mathbf{x_n}\} := x_0, x_1, \ldots, x_{N-1}$ into another sequence of complex numbers, $\{\mathbf{X_k}\} := X_0, X_1, \ldots, X_{N-1}$, which is defined by

$$X_k = \sum_{n=0}^{N-1} x_n \cdot e^{-\frac{i2\pi}{N}kn} = \sum_{n=0}^{N-1} x_n \cdot \left[\cos\left(\frac{2\pi}{N}kn \right) - i \cdot \sin\left(\frac{2\pi}{N}kn \right) \right].$$

- The inverse transform is given by:

$$x_n = \frac{1}{N} \sum_{k=0}^{N-1} X_k \cdot e^{i2\pi kn/N}$$

- The DFT is a linear transformation and thus can also be written in terms of a DFT matrix $\mathbf{F}$
- The transform is sometimes denoted by the symbol $\mathcal{F}$ or $\mathbf{F}$, as in $\mathbf{X} = \mathcal{F}\{\mathbf{x}\}$ or $\mathcal{F}(\mathbf{x})$ or $\mathbf{Fx}$.
- When scaled appropriately it becomes a unitary matrix $\mathcal{F}\mathcal{F}^{-1} = I$ and the X_k thus be viewed as coefficients of x in an orthonormal basis.

FFT: Fast Fourier Transform

$$\sum_{n=1}^{5} n \times cos(n \times \omega \times t), \quad \omega = 10 \times 2\pi$$

- Evaluating this definition directly requires $O\left(N^2\right)$ operations: there are N outputs X_k, and each output requires a sum of N terms. An FFT is any method to compute the same results in $O(N \log N)$ operations.
- The Fast Fourier Transform is computed on a uniformly sampled array and returns a uniformly sampled **spatial frequency decomposition**.

NFFT: Non-equispaced Fourier Transform

- The NFFT is a generalization of the FFT so that it works on a non-uniformly sampled signal and/or returns non-uniformly sampled spatial frequency decomposition.

Variance and covariance: a reminder

- Variance of a scalar-valued random variable X:
  $\sigma_X^2 = \text{var}(X) = E[(X - E[X])^2] = E[(X - E[X]) \cdot (X - E[X])]$

- Covariance between two scalar-valued random variables X and Y:

$$\text{cov}(X, Y) = E\left[(X - E[X])(Y - E[Y])\right] = E[XY] - E[X]E[Y]$$

- Covariance matrix of random vector $\boldsymbol{X}$:

$$\Sigma_{\boldsymbol{XX}} = \begin{bmatrix} E[(X_1 - E[X_1])(X_1 - E[X_1])] & E[(X_1 - E[X_1])(X_2 - E[X_2])] & \cdots & E[(X_1 - E[X_1])(X_n - E[X_n])] \\ E[(X_2 - E[X_2])(X_1 - E[X_1])] & E[(X_2 - E[X_2])(X_2 - E[X_2])] & \cdots & E[(X_2 - E[X_2])(X_n - E[X_n])] \\ \vdots & \vdots & \ddots & \vdots \\ E[(X_n - E[X_n])(X_1 - E[X_1])] & E[(X_n - E[X_n])(X_2 - E[X_2])] & \cdots & E[(X_n - E[X_n])(X_n - E[X_n])] \end{bmatrix}$$
$$= E[(\boldsymbol{X} - E[\boldsymbol{X}])^\top(\boldsymbol{X} - E[\boldsymbol{X}])]$$

- Cross-covariance matrix of random vectors $\boldsymbol{X}$ and $\boldsymbol{Y}$:
  $\Sigma_{\boldsymbol{XY}} = E[(\boldsymbol{X} - E[\boldsymbol{X}])^\top(\boldsymbol{Y} - E[\boldsymbol{Y}])]$

Linear regression

- $y_i = f(x_i) + \epsilon_i = \theta_1 + \theta_2 x_i + \epsilon_i$ where $\epsilon \sim \mathcal{N}(0, \sigma^2)$
- In matrix notation

$$
\boldsymbol{Y} = \begin{bmatrix} 1 & 1 & \ldots & 1 \\ x_1 & x_2 & \ldots & x_N \end{bmatrix}^\top \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} + \boldsymbol{\epsilon}
$$

$$
= \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix}^\top \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} + \boldsymbol{\epsilon} = \boldsymbol{X}^\top \boldsymbol{\theta} + \boldsymbol{\epsilon} \quad \text{where } \boldsymbol{\epsilon} \sim \mathcal{N}(\boldsymbol{0}, \boldsymbol{\Sigma})
$$

- Non-diagonal $\boldsymbol{\Sigma}$ will be used in the case the data points are covariant
- Likelihood of θ:

$$
\mathcal{L}(\boldsymbol{\theta}|\boldsymbol{Y}) = \Pr(\boldsymbol{Y}|\boldsymbol{\theta}) = \frac{1}{\sqrt{(2\pi)^N \det(\boldsymbol{\Sigma})}} e^{-\frac{1}{2}(\boldsymbol{Y}-\boldsymbol{X}^\top \boldsymbol{\theta})^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{Y}-\boldsymbol{X}^\top \boldsymbol{\theta})}
$$

- Likelihood of θ:

$$\mathcal{L}(\boldsymbol{\theta}|\boldsymbol{Y}) = \Pr(\boldsymbol{Y}|\boldsymbol{\theta}) = \frac{1}{\sqrt{(2\pi)^N \det(\boldsymbol{\Sigma})}} e^{-\frac{1}{2}(\boldsymbol{Y}-\boldsymbol{X}^{\top}\boldsymbol{\theta})^{\top}\boldsymbol{\Sigma}^{-1}(\boldsymbol{Y}-\boldsymbol{X}^{\top}\boldsymbol{\theta})}$$

- Maximum likelihood (take the log, derive with respect to θ) results in the **normal equation** giving the most likely $\boldsymbol{\theta}$:

$$\hat{\boldsymbol{\theta}} = (\boldsymbol{X}\boldsymbol{\Sigma}^{-1}\boldsymbol{X}^{\top})^{-1}\boldsymbol{X}\boldsymbol{\Sigma}^{-1}\boldsymbol{Y}$$

- Given $\hat{\boldsymbol{\theta}}$ and any new $\boldsymbol{X}_*$, we can now predict $\boldsymbol{Y}_*$, using the predictor **projection "hat" matrix $\boldsymbol{H}$** defined as:

$$\boldsymbol{Y}_* = (\boldsymbol{X}_*)^{\top}\hat{\boldsymbol{\theta}} = \underbrace{(\boldsymbol{X}_*)^{\top}(\boldsymbol{X}\boldsymbol{\Sigma}^{-1}\boldsymbol{X}^{\top})^{-1}\boldsymbol{X}\boldsymbol{\Sigma}^{-1}}_{\boldsymbol{H}}\boldsymbol{Y} = \boldsymbol{H}\boldsymbol{Y}$$

- All this is not fully Bayesian... What about $\Pr(\hat{\boldsymbol{\theta}})$, or $\Pr(\boldsymbol{Y}_*)$?

Marginal and Conditional Gaussians

- Important theorem, used when both the prior and posterior are normally distributed. If we have:

$$\Pr(\boldsymbol{x}) \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Lambda})$$
$$\Pr(\boldsymbol{y}|\boldsymbol{x}) \sim \mathcal{N}(\boldsymbol{A}\boldsymbol{x} + \boldsymbol{b}, \boldsymbol{\Sigma})$$

then we will have

$$\Pr(\boldsymbol{y}) \sim \mathcal{N}(\boldsymbol{A}\boldsymbol{\mu} + \boldsymbol{b}, \boldsymbol{\Sigma} + \boldsymbol{A}\boldsymbol{\Lambda}\boldsymbol{A}^\top)$$
$$\Pr(\boldsymbol{x}|\boldsymbol{y}) \sim \mathcal{N}((\boldsymbol{\Lambda}^{-1} + \boldsymbol{A}^\top\boldsymbol{\Sigma}^{-1}\boldsymbol{A})^{-1}\left(\boldsymbol{A}^\top\boldsymbol{\Sigma}^{-1}(\boldsymbol{y} - \boldsymbol{b}) + \boldsymbol{\Lambda}^{-1}\boldsymbol{\mu}\right), (\boldsymbol{\Lambda}^{-1} + \boldsymbol{A}^\top\boldsymbol{\Sigma}^{-1}\boldsymbol{A})^{-1})$$

- Reference: Bishop "Pattern Recognition and Machine Learning", eqs 2.113 to 2.117

Application to the Linear Regression case

- Somewhat confusing, but we have: $y \to Y$, $x \to \theta$, $A \to X^\top$, $b \to 0$, and $\Sigma \to \Sigma$
- We need to set $\Pr(\theta) \sim \mathcal{N}(\mu, \Lambda)$ with possibly the edge case of $\mu = 0$ and $\Lambda^{-1} \to 0$ to simulate a nearly uniform prior.
- We already have:

$$\Pr(Y|\theta, X) \sim \mathcal{N}(X^\top \theta, \Sigma)$$

- The posterior distribution for θ, whose mean is the MAP:

$$\Pr(\theta|Y, X) \sim \mathcal{N}((\Lambda^{-1} + X\Sigma^{-1}X^\top)^{-1}\left(X\Sigma^{-1}Y + \Lambda^{-1}\mu\right), (\Lambda^{-1} + X\Sigma^{-1}X^\top)^{-1})$$
$$= \mathcal{N}(\Gamma^{-1}X\Sigma^{-1}Y, \Gamma^{-1}) \text{ for common case } \mu = 0$$

- And the predictive distribution is (applying the previous slide again):

$$\Pr(Y_*|X_*, X, Y) = \int \Pr(Y_*|\theta, X_*, X, Y)\Pr(\theta|X, Y)d\theta$$
$$\sim \mathcal{N}(X_*^\top \Gamma^{-1}X\Sigma^{-1}Y, X_*^\top \Gamma^{-1}X_*)$$

Beyond Linear Regression: basis functions

- Before we had $\boldsymbol{Y} = \boldsymbol{X}^\top \boldsymbol{\theta}$, but now we'd prefer a more flexible scheme $\boldsymbol{Y} = \phi(\boldsymbol{X})^\top \boldsymbol{\theta}$ where ϕ represent a function basis such as polynomials $\phi(\boldsymbol{X}) = (1, x, x^2, x^3)^\top$. The matrix $\Phi(\boldsymbol{X})$ is the aggregation of columns $\phi(\boldsymbol{X})$ into a matrix.

- Amazingly, the same analysis works, so that the predictive distribution becomes:

$$\text{Pr}(\boldsymbol{Y}_*|\boldsymbol{X}_*, \boldsymbol{X}, \boldsymbol{Y}) \sim \mathcal{N}(\Phi(\boldsymbol{X}_*)^\top \boldsymbol{\Gamma}^{-1} \Phi(\boldsymbol{X}) \boldsymbol{\Sigma}^{-1} \boldsymbol{Y}, \Phi(\boldsymbol{X}_*)^\top \boldsymbol{\Gamma}^{-1} \Phi(\boldsymbol{X}_*))$$

where $\Gamma = \boldsymbol{\Lambda}^{-1} + \Phi(\boldsymbol{X})^\top \Sigma^{-1} \Phi(\boldsymbol{X})$

- This means we can find predictive distribution for basis functions. But what if we don't want to specify any functional form for these functions ?

Gaussian process

- A Gaussian process is a collection of random variables, any finite number of which have a joint Gaussian distribution.
- A Gaussian process is a distribution over functions, rather than over variables
- We define the mean function $m(\boldsymbol{x})$ and the covariance function $k(\boldsymbol{x}, \boldsymbol{x'})$ of a real process $f(\boldsymbol{x})$ and we note $\boldsymbol{f} \sim \mathcal{GP}(m(\boldsymbol{x}), k(\boldsymbol{x}, \boldsymbol{x'}))$ so that:

$$m(\boldsymbol{x}) = E[f(\boldsymbol{x})]$$
$$k(\boldsymbol{x}, \boldsymbol{x'}) = E\left[(f(\boldsymbol{x}) - m(\boldsymbol{x}))\left(f(\boldsymbol{x'}) - m(\boldsymbol{x'})\right)\right]$$

The squared exponential kernel

- The squared exponential is a classic choice for a first exposition to Gaussian processes

$$\mathbf{\Sigma} = cov(f(x), f(x')) = k(x, x') = \sigma_k^2 e^{-\frac{1}{2}\frac{(x-x')^2}{l^2}}$$

- It is parametrized by hyperparameter l, the characteristic length-scale of the process over which correlation is strong, and σ_k, the strength of correlation.
- Samples $\mathbf{Y} \sim \mathcal{N}(\mathbf{M}, K)$ can be generated by :
    1. computing $\mathbf{\Sigma}$ using the kernel expression over a given range for $\mathbf{X}$.
    2. computing the Cholesky decomposition $\mathbf{L}$ of $\mathbf{\Sigma} = \mathbf{L}\mathbf{L}^\top$
    3. computing $\mathbf{Y} = \mathbf{m} + \mathbf{L}\mathbf{u}$ where $\mathbf{u} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- Demo:

  https://distill.pub/2019/visual-exploration-gaussian-processes/

Predictions: posterior

- Typical example is data $\boldsymbol{Y} = f(\boldsymbol{X}) + \epsilon$, with $\epsilon \sim \mathcal{N}(\boldsymbol{0}, \sigma_n^2 \boldsymbol{I})$.
- $\text{cov}(y_i, y_j) = k(x_i, y_j) + \sigma_n^2 \delta_{ij} \rightarrow \text{cov}(\boldsymbol{Y}) = \boldsymbol{K}(\boldsymbol{X}, \boldsymbol{X}) + \sigma_n^2 \boldsymbol{I}$
- The joint distribution of the measurements and the predictions is:

$$\begin{bmatrix} \boldsymbol{Y} \\ \boldsymbol{f}_* \end{bmatrix} \sim \mathcal{N}\left(\boldsymbol{0}, \begin{bmatrix} \boldsymbol{K}(\boldsymbol{X}, \boldsymbol{X}) + \sigma_n^2 \boldsymbol{I} & \boldsymbol{K}(\boldsymbol{X}, \boldsymbol{X}_*) \\ \boldsymbol{K}(\boldsymbol{X}_*, \boldsymbol{X}) & \boldsymbol{K}(\boldsymbol{X}_*, \boldsymbol{X}_*) \end{bmatrix} \right)$$

- The predictive distribution is: $\text{Pr}(\boldsymbol{f}_* | \boldsymbol{X}, \boldsymbol{Y}, \boldsymbol{X}_*) \sim \mathcal{N}\left(\bar{\boldsymbol{f}}_*, \text{cov}(\boldsymbol{f}_*) \right)$, with

$$\bar{\boldsymbol{f}}_* = \boldsymbol{K}(\boldsymbol{X}_*, \boldsymbol{X}) \left(\boldsymbol{K}(\boldsymbol{X}, \boldsymbol{X}) + \sigma_n^2 \boldsymbol{I} \right)^{-1} \boldsymbol{Y}$$

$$\text{cov}(\boldsymbol{f}_*) = \boldsymbol{K}(\boldsymbol{X}_*, \boldsymbol{X}_*) \left(\boldsymbol{K}(\boldsymbol{X}, \boldsymbol{X}) + \sigma_n^2 \boldsymbol{I} \right)^{-1} \boldsymbol{K}(\boldsymbol{X}_*, \boldsymbol{X})$$

Marginalization of kernel parameters

- The prior $\Pr(\boldsymbol{f}|\boldsymbol{X}) \sim \mathcal{N}(\boldsymbol{0}, \boldsymbol{K})$ and likelihood $\Pr(\boldsymbol{Y}|\boldsymbol{f}) \sim \mathcal{N}(\boldsymbol{f}(\boldsymbol{X}), \sigma_n^2 \boldsymbol{I})$ give the marginal likelihood over the function values $\Pr(\boldsymbol{Y}|\boldsymbol{X}) \sim \mathcal{N}(\boldsymbol{0}, \boldsymbol{K} + \sigma_n^2 \boldsymbol{I})$

- For kernels with hyperparameters (e.g. σ_k and l for the squared exponential), one can write the expression for the log-marginal likelihood as a function of the hyperparameters, then optimize it.

- This allows us to find the best parameters supported by the data, and in turn to refine our future predictions.

- "Gaussian Processes for Machine Learning", Rasmussen & Williams, PDF downloadable here:
  http://www.gaussianprocess.org/gpml/
- "Pattern Recognition and Machine Learning", Bishop, with examples by contributors in Matlab http://prml.github.io/ or Python
  https://github.com/ctgk/PRML